

Adrian Cole | Data Scientist

Specializing in Machine Learning and Predictive Analytics

456 Mission Street, San Francisco, California, United States, 94110

☎ (415) 555-0142 • ✉ adrian.cole@example.com • 🌐 <https://adriancole.dev>

🐙 [@adriancole](https://github.com/adriancole) • 🌐 [@adriancole-data](https://www.linkedin.com/company/adriancole-data)

Basics

Data scientist with 6+ years of experience building machine learning systems that turn messy data into measurable business impact. Strong background in statistics, causal inference, and natural language processing, with deep expertise across the full model lifecycle—from problem framing and data pipelines to deployment and monitoring. Passionate about creating interpretable, reliable models and translating data insights into clear recommendations for product and leadership teams.

Education

University of California, Berkeley

<https://www.ischool.berkeley.edu/>

Master, Information and Data Science, Score: 3.9

Aug 2016–May 2018

- Concentrated on machine learning, causal inference, and large-scale data systems
- Capstone project: built a real-time churn prediction pipeline for a telecommunications partner
- Graduate research assistant supporting a study on algorithmic fairness in lending decisions

Courses: Applied Machine Learning, Statistics for Data Science, Natural Language Processing, Big Data Systems, Experimentation and Causal Inference, Data Visualization

University of Washington

<https://www.stat.washington.edu/>

Bachelor, Statistics, Score: 3.8

Sep 2012–Jun 2016

- Built strong foundations in statistical theory and computational methods
- Undergraduate research assistant in the Statistical Learning Lab

Courses: Statistical Inference, Regression Models, Probability Theory, Computational Statistics, Data Mining

Work

Lumina Analytics

<https://www.lumina-analytics.com/>

Senior Data Scientist

Mar 2021–Present

- Lead development of predictive models for customer lifetime value across a portfolio of SaaS clients, improving average revenue forecast accuracy by 18%
- Architect feature store and model monitoring pipelines using Python and MLflow, reducing model retraining time from days to hours
- Partner with product and engineering teams to design A/B tests and causal inference studies for new feature adoption initiatives
- Mentor a team of four data scientists and establish best practices for code review, reproducibility, and documentation

Keywords: Python, MLflow, Feature Store, Causal Inference, Mentorship

HealthInsight

<https://www.healthinsight.com/>

Data Scientist

Jul 2018–Feb 2021

- Built risk stratification models using electronic health records and claims data, enabling early intervention for high-risk patients
- Designed and implemented an NLP pipeline to extract clinical concepts from unstructured notes, achieving 87% F1 score
- Created interactive Tableau dashboards used by clinical leadership to track quality metrics and reduce hospital readmission rates by 12%

Keywords: Healthcare, NLP, Tableau, Electronic Health Records

Languages

English: Native or Bilingual Proficiency

Keywords: IELTS 8.5

Spanish: Full Professional Proficiency

Keywords: DELE C1

Skills

Machine Learning: Expert

Keywords: Python, scikit-learn, XGBoost, TensorFlow, PyTorch

Statistical Modeling: Advanced

Keywords: R, Bayesian Methods, Time Series, Survival Analysis, Experiment Design

Data Engineering: Advanced

Keywords: SQL, Spark, dbt, Airflow, AWS, Snowflake

Natural Language Processing: Advanced

Keywords: Hugging Face, spaCy, Transformers, NER, Topic Modeling

Data Visualization: Intermediate

Keywords: Tableau, matplotlib, Plotly, ggplot2

Awards

Outstanding Contribution in Data Science

Lumina Analytics

Dec 2023

Recognized for leading the development of a company-wide model monitoring platform that reduced unplanned model retraining incidents by 40%.

Certificates

AWS Certified Machine Learning - Specialty <https://aws.amazon.com/certification/certified-machine-learning-specialty/>

Amazon Web Services

Nov 2021

TensorFlow Developer Certificate

<https://www.tensorflow.org/certificate>

Google

Apr 2020

Publications

Proceedings of the ACM SIGKDD Conference on Knowledge Discovery and Data Mining <https://doi.org/10.1145/3523227.3547400>

A Scalable Framework for Real-Time Customer Churn Prediction

Aug 2022

- Presented a production-ready churn prediction framework combining gradient-boosted trees with online feature stores
- Demonstrated a 23% relative improvement in AUC versus batch baselines across three industrial datasets

Projects

DocQuery

<https://github.com/adriancole/docquery>

An NLP-powered search and question-answering tool for unstructured clinical documents.

Jan 2023–Jun 2023

- Developed a retrieval-augmented system using sentence embeddings and a fine-tuned transformer model to answer clinician questions from EHR notes
- Deployed as a REST API with Docker, reducing average answer lookup time from 20 minutes to under 5 seconds

Keywords: NLP, Transformers, Q&A, Docker, FastAPI

AdScore

<https://github.com/adriancole/adscore>

Click-through rate prediction model for online advertising campaigns.

Aug 2020–Dec 2020

- Engineered features from impression logs and built a distributed training pipeline in PySpark
- Achieved a 15% lift in offline AUC and enabled real-time bidding through model export to ONNX

Keywords: CTR, PySpark, Feature Engineering, ONNX

Interests

Open Source: Contributing to scikit-learn, Jupyter, Streamlit

Hiking: Yosemite, Pacific Crest Trail

Chess: Tournament Play, Puzzle Solving

Volunteer

Bay Area Data Corps

<https://www.bayareadatacorps.org/>

Data Science Volunteer

Mar 2022–Dec 2023

- Partnered with local nonprofits to analyze service utilization data and guide resource allocation decisions
- Built reusable notebook templates to standardize weekly reporting and make findings accessible to nontechnical staff